

Highlights

- A size-invariant deep RL policy schedules interval job shops of any size
- Trained on one size class, it beats every priority rule on all 70 instances
- Thirty artifacts per paradigm, compared in one harness at matched budgets
- Which learner leads is decided by the inference budget and the size class
- A width penalty acts on the policy's selection, on the rule's representation