

Highlights

- A size-invariant deep RL policy schedules interval job shops of any size
- Trained on one size class, it beats every priority rule on all 70 instances
- Thirty artifacts per paradigm compared in one harness at one constructive pass
- Which learner leads is decided by the inference budget and the size class
- A width penalty narrows policy schedules mainly by selection, GP by structure