

Supplementary material

Genetic Programming Hyper-Heuristics for the Job Shop Scheduling Problem with Interval Durations: Robustness-Aware Interpretable Rules that Generalize Across Instance Sizes

Sections, tables, figures and equations without the prefix S refer to the article.

S1 Interval handling and the decoder

Two aspects of the comparison with the constructive baselines are examined here: how each baseline reduces an interval to a comparable quantity, and how much of the margin of the evolved rules belongs to the priority rule rather than to the schedule generation scheme.

The baselines rank candidates by the lexicographic order of Eq. (8) (Section 5.2). To check that this choice does not handicap them, each interval-valued rule was re-run ranking by a single number drawn from the same interval: its lower endpoint, its midpoint or its upper endpoint. Over the 70 instances SPT spans 633.1 to 642.5 across those three conventions, LPT 703.1 to 704.4, MWKR 62.9 to 64.6 and EST 42.7 to 45.1. The lexicographic values of Table 4 are better than all three for SPT and EST, and within 1.9 points of the best for LPT and MWKR. For EST and MWKR, the rules that account for the state of the shop, no convention moves the result by more than 2.4 points, and for SPT and LPT, whose levels exceed 600, by at most 9.4; the best plain rule and the evolved one are 23.3 points apart. How the intervals are collapsed is therefore not what the comparison measures.

The decoder is a more substantial question, since the strongest baseline, G&T-MWKR, does not use the same one: it restricts each decision to Giffler and Thompson’s conflict set and so produces active schedules, whereas the evolved rules choose from the whole eligible set and produce semi-active ones. Running the evolved rules as the tie-break *inside* the conflict set separates the two effects. Table S1 gives the four combinations. Holding the decoder fixed, the evolved rules outperform MWKR under either one: 18.99 against 64.7 in the semi-active decoder and 23.18 against 29.5 in G&T’s. Holding the rule fixed, each prefers a different decoder: the evolved rules lose 4.2 points when moved into the conflict set, while MWKR gains 35.2 by entering it. Each method in Table 4 is thus reported with the decoder that suits it, and the margin survives the other choice.

Table S1: RE (%) over the 70 instances for the two decoders and the two priority rules, one deterministic pass. The GP entries are the mean over the 30 evolved rules.

Priority rule	Semi-active	Giffler–Thompson
MWKR	64.7	29.5
Evolved, mean of 30	18.99 $\pm$ 1.33	23.18 $\pm$ 7.47
Evolved, featured	17.71	18.15

Two details of the table deserve mention. The featured rule is the one case where the decoder barely matters, 17.71 against 18.15, a difference of 0.44 points that a paired test over the 70 instances does not separate from zero ($p = 0.32$); the loss is concentrated in the rest of the arm, whose standard deviation grows from 1.33 to 7.47 because a few rules deteriorate markedly once their candidate set is restricted. This is what should be expected of rules evolved against a particular generation scheme: the conflict set removes candidates their fitness was never measured on. And SPT, which is the worst entry of Table 4 at 627.2, reaches 70.6 inside the conflict set, as G&T-SPT, which is the same phenomenon seen from the other side.

S2 Robustness under executional uncertainty: further analyses

This section extends the robustness comparison of Section 7.6 of the article, whose Table 8 the paragraphs below refer to. It gives the per-instance distributions under widened intervals, reports the deviation in time units, $|\Delta|$, repeats the comparison at the level of the 30 rules of each arm, changes the law of the realizations, and describes the worst 5% of executions.

In time units the featured rule deviates by 12.30, against 14.64 for EST, and the difference is significant ($z = -5.89$, $p < 0.001$, $|r| = 0.81$). Against G&T-MWKR the ordering is unchanged and the gap wider (12.30 against 14.94, $z = -6.83$, $|r| = 0.94$). The featured rule and its ablated counterpart remain indistinguishable (12.30 against 12.25, $z = 0.61$, $p = 0.55$).

The robust rule also deviates less than its own ablated counterpart, by 1.78 time units ($z = -5.17$, $|r| = 0.71$).

The comparisons above select one representative rule per arm; at the level of the method the ordering is the same. Recomputing $\bar{\varepsilon}$ at the nominal width for all 30 rules of each arm and comparing the arms as independent samples, the makespan arm averages 5.67 ± 0.16 against 5.77 ± 0.14 for its ablation ($z = -2.40$, $p < 0.05$ after Holm's correction over the three arm-level tests, $|r| = 0.36$). The robust $\lambda=1$ arm averages 5.50 ± 0.30 against 5.68 ± 0.11 for its own ablation ($z = -2.91$, $p < 0.05$, $|r| = 0.44$), and it improves on the makespan arm with the same terminals ($z = -2.68$, $p < 0.05$, $|r| = 0.40$).

Two things follow. The width terminals reduce $\bar{\varepsilon}$ by 1.7% even under the makespan objective, an effect too small to be detected with a single pair of rules. And the arm-level reductions are modest, three percent or less, against the fifth separating the representative rules, so selecting the narrowest rule of the robust arm on the validation set amplifies an effect that is consistent but small at the level of the method.

The estimate of the article draws each realization uniformly on its interval, which is one choice among several. Repeating the measurement with a symmetric triangular law, with a law skewed towards the upper endpoint, and with the single worst-case realization in which every duration takes its upper bound, leaves the ordering and every significant contrast intact. Averaged over the 70 instances, the robust $\lambda=1$ rule improves on the makespan rule by 1.22 time units under the triangular law, 12.16 under the skewed one and 29.01 in the worst case, all with $p < 0.001$. The featured rule and its ablated counterpart remain indistinguishable under all three. The deviations themselves grow with the mass the law places near the upper endpoint, from 8.95 time units for the rule under the triangular law to 57.15 under the skewed one and 138.35 in the worst case, but the comparison between methods does not depend on which law is used.

The measures above are averages over realizations, whereas a planner exposed to late completions is concerned with the worst ones. The conditional value-at-risk of Section 5.4, computed on the same $K=1000$ realizations, describes the 5% worst executions of each schedule and separates two questions. The first is how far those executions exceed the prediction. Under the uniform law the featured rule overruns its expected makespan by 33.46 time units on average over its worst 5% of executions, against 38.87 for G&T-MWKR and 40.81 for EST ($z = -5.04$ and $z = -5.61$, $p < 0.001$), and its ablated counterpart is indistinguishable (33.55, $z = -0.05$). The robust objective reduces the tail overrun in about the same proportion as the mean deviation: to 27.46 at $\lambda=1$, 17.9% less than the makespan rule ($z = -5.14$, $p < 0.001$, $|r| = 0.71$) and 5.68 less than its own ablation ($z = -5.32$, $|r| = 0.73$), and to 26.43 at $\lambda=4$. At matched a-priori quality, the robust $\lambda=4$ rule overruns by 26.43 against 38.87 for G&T-MWKR, 32.0% less ($z = -7.00$, $|r| = 0.96$). The second question is how long the worst executions are in absolute terms, and there the ordering follows expected makespan rather than robustness. Expressed as RE, the tail of the featured rule lies at 19.53% and those of the robust rules at 26.46% and 29.73%, higher on 68 and on 70 of the 70 instances, respectively. The robust objective therefore makes the prediction more reliable in the tail as well as on average, but it does not shorten the worst executions, whose length follows the expected makespan given up for that reliability. Against

G&T-MWKR, whose expected makespan is close to its own, the $\lambda=4$ rule has a slightly shorter tail (29.73 against 31.59, $z = -2.41$, $p = 0.016$, $|r| = 0.33$). All of these contrasts keep their sign under the triangular and the skewed laws, under which the tail overrun at $\lambda=1$ falls by 3.23 and 14.89 time units against the makespan rule, and all keep their significance except the shorter tail of the $\lambda=4$ rule against G&T-MWKR, which is marginal under the triangular law ($p = 0.044$).




Figure S1: Per-instance $\bar{\varepsilon}$ by method at nominal interval widths and widened by +20% and +40%; the robust $\lambda=1$ entries are the rules of Table 8.

Figure S1 shows the per-instance $\bar{\varepsilon}$ distributions. The ordering is stable as uncertainty grows, and the advantage of the robust rule widens with it: at +40% its deviation is 6.08 against 7.57 for the makespan rule and 8.58 for G&T-MWKR.

S3 Sensitivity to the training set

The rules of the article are evolved on four instances, TA11–TA14, fixed before any result was inspected. Five further campaigns repeat the protocol unchanged over other training sets. Each runs thirty independent evolutions with the same terminal set and the same irace configuration (Section 5.3 of the article). Three campaigns use alternative quadruples of the 20×15 class: TA15–TA18, TA17–TA20, and the even-numbered TA12, TA14, TA16 and TA18. The other two change the size of the training set, to a pair (TA11–TA12) and an octuple (TA11–TA18). All six campaigns are evaluated on the 60 instances of the other size classes, which none of them trains on. Each campaign is summarized by the mean over its thirty rules, with no selection step, so the validation set plays no part in the comparison.

Table S2 gives the outcome. The six means span 0.66 points, from 18.95% to 19.61%. The standard deviation between seeds within a campaign, 0.93–1.35, is larger than that span: which instances are trained on matters less than which seed the evolution starts from. Each campaign is compared with the reference as an independent sample of rules (Section 5.4 of the article). No campaign differs significantly from TA11–TA14, even before Holm’s correction over the five contrasts: the smallest p is 0.06, and 0.31 after correction. The size of the training set has an

Table S2: RE (%) over the 60 instances that no campaign trains on, by training set: mean $\pm$ sd across 30 independently evolved rules, best rule, difference of means from the reference, and the two-sided Mann–Whitney contrast against the reference.

Training set	$ T $	Mean $\pm$ sd	Best	Δ	z	p
TA11–TA14 (reference)	4	19.15 ± 1.35	17.89	—	—	—
TA15–TA18	4	19.46 ± 1.24	17.57	+0.31	+1.40	0.162
TA17–TA20	4	19.53 ± 1.07	17.76	+0.37	+1.85	0.066
TA12, 14, 16, 18	4	19.07 ± 0.93	17.49	−0.08	+0.38	0.706
TA11–TA12	2	19.61 ± 1.23	17.53	+0.46	+1.88	0.061
TA11–TA18	8	18.95 ± 1.03	17.38	−0.20	−0.21	0.842

effect of the expected sign and small magnitude. Halving it to two instances costs 0.46 points and doubling it to eight gains 0.20, both within the variation between seeds. The best rule of each campaign ranges from 17.38% to 17.89%, and the reference quadruple sits at the unfavorable end of that range.

What the training set does change is the shape of the rules. SLACK leads the terminal count under every training set, with 21% to 27% of the leaves. The width terminal WKRW loses ground: it falls from 12.0% under TA11–TA14 to 5.8–11.8% under the other sets. EST rises from 9.7% to 10.8–14.4%. Different instances select different tie-breakers around the same SLACK-led core, and reach the same performance with them. The featured rule of the article is thus one good solution among several, not the only one the method can produce.

S4 Asymmetric intervals

All the intervals used in the article are symmetric about the crisp duration. To test whether the contribution of the width terminals depends on that construction, the four arms of Sections 7.4 and 7.5 were re-evolved on the asymmetric counterparts of TA11–TA14 described in Section 5.1, with the same configuration and fifteen seeds per arm rather than thirty, a reduction imposed by the computational cost. Table S3 reports them over all 70 asymmetric instances, with the deviation of the executed makespan estimated from 400 realizations of the fixed processing order. Levels of RE do not compare across the two benchmarks, since the right skew raises the expected makespan while the crisp lower bound does not move and every arm sits some six points higher; only the contrast between arms within a benchmark is meaningful.

Table S3: The four arms re-evolved on asymmetric intervals, over the 70 asymmetric instances: mean $\pm$ sd across 15 seeds of RE, relative width of the predicted interval and absolute deviation of the executed makespan.

Arm	RE (%)	Width (%)	$ \Delta $
Makespan, full	25.08 ± 1.42	13.23 ± 0.72	13.36 ± 0.64
Makespan, no width	23.90 ± 1.18	13.08 ± 0.20	13.06 ± 0.26
Robust $\lambda=1$, full	25.93 ± 1.18	12.67 ± 0.20	12.75 ± 0.27
Robust $\lambda=1$, no width	23.89 ± 1.06	12.97 ± 0.18	13.03 ± 0.16

The main result is reproduced. Under the robust objective the width terminals narrow the predicted interval from 12.97 to 12.67 ($z = -3.59$, $|r| = 0.77$) and reduce the deviation of the executed makespan from 13.03 to 12.75 ($z = -2.80$, $|r| = 0.60$), at a cost of 2.04 points of RE ($z = 3.59$). The robust arm also improves on the makespan arm carrying the same terminals, by 0.55 points of width and 0.61 time units of deviation ($z = -3.59$, $|r| = 0.77$ for the latter). These five contrasts survive Holm’s correction over the nine Mann–Whitney tests of the table, the largest adjusted p among them being 0.0273. The robustness that the width terminals

make attainable, and the objective under which they attain it, therefore do not depend on the symmetric generation scheme.

Under the makespan objective the results are weaker than on the symmetric benchmark. There, the width terminals made the predicted interval 0.20 points narrower even though that objective does not reward it, a significant difference (Section 7.4). Here they do not (13.23 against 13.08, $p = 0.59$), nor do they reduce the executed deviation (13.36 against 13.06, $p = 0.17$). They cost 1.18 points of RE, a difference that does not survive the correction ($p = 0.14$ adjusted). That the width terminals do not lower the expected makespan holds on both benchmarks, and the remaining results are consistent with Section 7.5: where the objective does not price width, the smaller terminal set suffices.

S5 Per-instance results

Table S4 lists, for each of the 70 interval Taillard instances, the crisp lower bound and the single-pass RE (%) of the featured rule and of the two strongest baselines: EST, the best of the plain dispatching rules, and G&T-MWKR, the best baseline overall.

Table S4: Per-instance RE (%) of the featured rule (GP) and of the two strongest baselines, EST and G&T-MWKR (single deterministic pass). LB is the crisp Taillard lower bound.

Inst.	LB	EST	G&T	GP	Inst.	LB	EST	G&T	GP
TA1	1231	49.9	26.1	19.3	TA36	1819	41.1	42.8	23.8
TA2	1244	24.9	27.8	14.2	TA37	1771	43.4	35.6	18.4
TA3	1218	33.0	33.2	14.1	TA38	1673	45.0	28.0	20.1
TA4	1175	32.0	27.7	13.0	TA39	1795	46.6	28.9	22.4
TA5	1224	27.2	33.0	14.7	TA40	1651	46.9	42.3	26.1
TA6	1238	26.8	18.2	13.7	TA41	1906	44.6	45.0	33.4
TA7	1227	30.7	26.1	17.7	TA42	1884	71.1	38.1	24.3
TA8	1217	23.7	25.0	14.8	TA43	1809	68.7	39.8	22.3
TA9	1274	37.7	28.2	20.0	TA44	1948	52.1	34.7	25.8
TA10	1241	37.0	22.1	15.4	TA45	1997	50.2	33.7	14.9
TA11	1357	56.2	30.7	17.5	TA46	1957	60.8	31.6	22.4
TA12	1367	48.8	34.6	17.6	TA47	1807	46.8	30.1	25.6
TA13	1342	60.2	26.9	13.3	TA48	1912	59.1	33.7	21.1
TA14	1345	39.5	29.7	13.2	TA49	1931	51.0	35.3	20.1
TA15	1339	48.1	30.6	16.7	TA50	1833	61.4	46.2	31.3
TA16	1360	45.2	34.0	15.6	TA51	2760	33.8	34.7	23.2
TA17	1462	58.9	17.7	13.9	TA52	2756	38.0	26.5	12.7
TA18	1377	43.8	31.8	23.6	TA53	2717	34.1	18.3	11.8
TA19	1332	45.8	29.3	20.0	TA54	2839	20.9	16.1	8.2
TA20	1348	45.7	24.1	15.1	TA55	2679	31.7	29.2	15.5
TA21	1642	33.4	28.0	20.2	TA56	2781	34.8	19.8	10.9
TA22	1561	44.2	30.8	17.4	TA57	2943	32.9	20.0	9.6
TA23	1518	42.8	35.3	17.8	TA58	2885	37.4	26.3	13.6
TA24	1644	56.3	23.8	14.1	TA59	2655	26.6	28.2	16.7
TA25	1558	45.8	33.8	16.6	TA60	2723	41.7	19.8	15.5
TA26	1591	42.8	31.2	26.9	TA61	2868	45.7	25.4	15.4
TA27	1652	50.0	31.1	21.4	TA62	2869	43.0	27.6	17.2
TA28	1603	36.8	26.8	12.4	TA63	2755	38.8	23.8	13.4
TA29	1583	28.8	32.0	17.2	TA64	2702	38.2	28.4	12.2
TA30	1528	34.2	36.1	17.0	TA65	2725	42.9	30.0	20.5
TA31	1764	30.7	34.4	24.5	TA66	2845	41.8	28.0	11.9
TA32	1774	42.7	33.3	19.0	TA67	2825	31.7	26.1	15.2
TA33	1788	51.3	34.8	27.5	TA68	2784	50.1	19.1	10.6
TA34	1828	45.9	31.8	20.7	TA69	3071	35.3	22.8	12.9
TA35	2007	30.3	24.1	8.4	TA70	2995	42.9	23.4	16.7

S6 The seeded genetic algorithm over time

Table S5 compares the genetic algorithm seeded with the rule’s permutation against the unseeded one at several budgets on the four instance sets of Section 6.5 of the article. The seed gives a significant advantage on every set after one second per instance, for longer the larger the instance, and none from 100 s per instance on. The tests at different budgets are not corrected

for multiplicity; with 10 instances the smallest attainable two-sided p is 0.002, the value repeated in the first rows.

Table S5: Seeded against unseeded genetic algorithm: instances on which the seeded one is better and, in parentheses, the two-sided paired Wilcoxon p , with the instance as unit. A dash marks a budget beyond the horizon of the set.

Seconds	Classical (12)	15×15 (10)	30×15 (10)	50×15 (10)
1	11 (0.001)	10 (0.002)	10 (0.002)	10 (0.002)
3	8 (0.064)	10 (0.002)	10 (0.002)	10 (0.002)
10	5 (0.970)	6 (0.625)	9 (0.004)	10 (0.002)
30	5 (0.970)	6 (0.625)	5 (0.846)	9 (0.006)
100	6 (0.791)	5 (1.000)	4 (0.432)	8 (0.064)
300	6 (0.791)	5 (0.846)	5 (0.846)	6 (0.432)
1000	—	—	4 (0.846)	6 (0.625)
End of horizon	6 (0.791)	3 (0.625)	4 (0.375)	7 (0.432)

S7 Ablation results

Table S6 gives the means of the arms shown in Figure 8 of the article, and the Mann–Whitney tests between them.

Table S6: Means of the ablation arms (30 evolutions each) and Mann–Whitney tests between them: the terminal ablation under two objectives, and the crisp midpoint control against the interval-trained arm above it. *Width* is the relative width of the predicted makespan interval.

Fitness objective	Terminal set	RE (%)	Width (%)
Makespan $E[C_{\max}]$	full	18.99 ± 1.33	12.41 ± 0.30
	without widths	18.66 ± 1.08	12.62 ± 0.23
	<i>test</i>	<i>n.s.</i> ($z = 0.67$)	$z = -2.96, p < 0.01$
Crisp midpoint control	without widths	18.63 ± 0.99	12.55 ± 0.17
	<i>test</i>	<i>n.s.</i> ($z = -0.12$)	<i>n.s.</i> ($z = -1.26$)
Robust, Eq. (13)	full	19.62 ± 1.52	12.04 ± 0.75
	without widths	18.42 ± 0.87	12.47 ± 0.20
	<i>test</i>	$z = 3.59, p < 0.001$	$z = -2.99, p < 0.01$